

Binary Search Trees and AVL Trees

Tufts University

Binary Search Trees - warmup questions


Q1: Create a BST by inserting the following keys in order: 7, 10, 5, 0, 9, 12, 15, 1, 4, 8, 11

Q2: What's the maximum height of a BST containing n elements?

Q3: What's the minimum height of a BST containing n elements?

Binary Search Trees

Container DS for fast insertion/deletions/search operations.




- ▶ Main structure is a **binary tree**
- ▶ Each node stores an object with a unique key
- ▶ **Main invariant:** objects in the right subtree of node n have smaller key than the object in n .

Similarly, left subtree stores objects with larger keys

Binary Search Trees

Container DS for fast insertion/deletions/search operations.




- ▶ Main structure is a **binary tree**
- ▶ Each node stores an object with a unique key

Main invariant: objects in the right subtree of node n have smaller key than the object in n .


Similarly, left subtree stores objects with larger keys

BST Operations



- ▶ Search
- ▶ Insert
- ▶ Predecessor/Successor
- ▶ Find minimum/maximum
- ▶ Delete
- ▶ (sorted) Print
- ▶ ...

Search/Insert



SEARCH(target)

$c \leftarrow \text{root}$

While ($c \neq \text{NIL}$)

 If $c.\text{key} = \text{target} \Rightarrow \text{Return true}$


 Else if $c.\text{key} < \text{target} \Rightarrow c \leftarrow c.\text{left}$

 Else $c \leftarrow c.\text{right}$

Return false

Exercise: modify code for an Insertion instead

Predecessor/Successor




SUCCESSOR(x): smallest number in T that is $\geq x$

- ▶ SUCCESSOR(1)=1
- ▶ SUCCESSOR(2)=4
- ▶ SUCCESSOR(16)= ∞
- ▶ ...

Strict Successor: smallest number in T that is $> x$

Further discussion




Let's find an algorithm for `STRICTSUCCESSOR`

- ▶ Prove that if y is the strict successor of x , then either y is an ancestor of x or x is an ancestor of y
- ▶ Give pseudocode for `STRICTSUCCESSOR`.
- ▶ Prove that your algorithm is correct. The above result may be helpful.

What would you change to compute `SUCCESSOR`?

What about `PREDECESSOR`/`STRICTPREDECESSOR`?

Finding Smallest/Largest element




Maximum right \Rightarrow right $\Rightarrow \dots$ until no right child

Minimum left \Rightarrow left $\Rightarrow \dots$ until no left child

Exercise: Prove that strategy works!

Deletions



Exercise! Things to look out for

Code depends on degree!


- ▶ Deleting a leaf is easy (just remove)
- ▶ Degree one → do a bypass
- ▶ Degree two → replace and recursively delete

Sorted print

Just do an **In-Order Traversal** of T

INORDER(x)

- ▶ if $x \neq \text{NIL}$ then
 - ▶ INORDER($x.\text{left}$)
 - ▶ Print(x)
 - ▶ INORDER($x.\text{right}$)



Exercise: What happens if we do preorder/postorder traversal?

Runtime?

What is the runtime of each of these operations?

- ▶ Search $\Theta(h)$
- ▶ Insert $\Theta(h)$
- ▶ Predecessor/Successor $\Theta(h)$
- ▶ Find minimum/maximum $\Theta(h)$
- ▶ Delete $\Theta(h)$
- ▶ (sorted) Print $\Theta(n)$
- ▶ ...

Where h is height of the tree $h = n$ in the worst case

Fun interlude: building a BST

- ▶ Remember heaps?

Single insertion in $\Theta(\log n)$

Insert n bottom-up in $\Theta(n)$

- ▶ Can we do something similar with BSTs? **Yes!**

RANDOMIZED-BST-BUILD(A)

Step 1: Randomly shuffle input array A

Step 2: For i from 1 to n insert $A[i]$ in BST

Runtime?

Example: Let's insert 6, 23, 12, 4, 20, 1, ...

Same as QUICKSORT!

First pivot = root of the BST

Second pivots = children of root

Etc ...

Correspondence between comparisons of both algorithms

Corollary: the expected height of the tree is $\Theta(\log n)$

⇒ n insertions will need $\Theta(n \log n)$ expected time

Balancing BSTs

Runtime of most operations in BSTs is $\Theta(h)$

If we insert all nodes at once we can have $h \approx \log n$

Can we maintain balance in other situations? **Yes!**


- ▶ 2-3 trees
- ▶ AA trees
- ▶ AVL trees \Leftarrow **this class**
- ▶ B-tree
- ▶ Red-black tree
- ▶ and more!

Introducing AVL trees


Adelson-Velsky and Landis trees (AVL trees for short) are BST with one more invariant:

- ▶ For any node, the difference in heights of subtrees is ≤ 1 .

Is this tree an AVL tree? **No!**



AVL trees - Height



Lemma: the height of an AVL tree is $\Theta(\log n)$.

Corollary: most operations need $\Theta(\log n)$ time.

Proof: next lecture.

Exercise: prove **Leaf Ratio Lemma**: If x and y are leaves then $\frac{d(y)}{d(x)} \leq 2$

$d(x)$ denotes the number of edges from the root to x .


What changes in AVL trees?

Compared to BST, what do we need to change in ...

- ▶ Search
- ▶ Insert ⇐
- ▶ Predecessor/Successor
- ▶ Find minimum/maximum
- ▶ Delete ⇐
- ▶ (sorted) Print
- ▶ ...

Inserting in AVL tree

What changes in an insertion?



Ex:

Insert(6)


Insert(3)

Insert(2)

After insertion we may have to rebalance

Rotations


ROTATE(Node, direction)



- ▶ Rotation preserves BST properties because $X \leq A \leq Y \leq B \leq Z$
- ▶ Let's use it to preserve balance invariant in AVL trees

Insertion in AVL trees


Back to INSERT(3)



5 is unbalanced (left child too deep) \Rightarrow Right rotate 5

AVL trees - Insert

INSERT(3)



5 is unbalanced (left child too deep) \Rightarrow Right rotate 5

General rule


After inserting node N travel back to root checking balance.

- ▶ If all nodes balanced \Rightarrow done!
- ▶ Otherwise, stop at first (i.e. deepest) unbalanced ancestor N
- ▶ Wlog, assume N is in the left subtree of B . Two cases:


N in left subtree of $A = B.left$

N in right subtree of $A = B.left$

CASE 1




CASE 2



AVL trees - Insert: Case 1


CASE 1



Done! No more height changes!

AVL trees - Insert: Case 2

CASE 2




Case 2 cannot be fixed with 1 rotation 😞

AVL trees - Insert: Case 2 FIXED


To fix Case 2 we “zoom into” $C == A.right$.

CASE 2A



How can we rebalance this tree? Hint: use 2 rotations.

AVL trees - Insert: Case 2A




No more height changes after rotation

Exercise: what if Y_1 is empty?

WOW!!!

What about deletions?



- ▶ Perform standard BST-Delete. Check upwards for unbalance
- ▶ As before, only need to check ancestors
- ▶ Rotations can be more complicated than Insert
(up to $\Theta(\log n)$ rotations)